

Children's social referencing reflects sensitivity to graded uncertainty

Emily Hembacher

ehembach@stanford.edu
Department of Psychology
Stanford University

Benjamin deMayo

bedemayo@stanford.edu
Department of Psychology
Stanford University

Michael C. Frank

mcfrank@stanford.edu
Department of Psychology
Stanford University

Abstract

Children likely rely on uncertainty monitoring to coordinate active learning behaviors. However, we still know little about children's ability to monitor epistemic uncertainty, given contradictory findings across tasks. We examined a spontaneous behavioral reaction to uncertainty, social referencing, during a word learning task among preschoolers. Children ages 2-5 were asked to place a target object in a bucket after hearing the experimenter produce a label for the target. We manipulated referential ambiguity through the number of objects present and their familiarity: in Experiment 1, when there were two unfamiliar objects and an unfamiliar label, the referent was unclear; when there were two familiar objects, or only one novel or familiar object, the referent was known or could be inferred. In Experiment 2, there were either two novel objects, two familiar objects, or one familiar and one novel object, in which case the referent could be inferred using the mutual exclusivity principle. Across Experiments 1 and 2, children looked up to the experimenter more often while executing a decision about which object to place in the bucket when there were two novel objects and thus referential ambiguity, and they did so while planning their decision as well in Experiment 1. In Experiment 2, children also referenced the experimenter more on mutual exclusivity trials, but only when the experimenter did not gaze at the object as she labeled it, suggesting that children's social referencing is sensitive to graded uncertainty.

Keywords: social referencing; help seeking; word learning; uncertainty.

Human learning can be characterized as a problem of detecting and reducing epistemic uncertainty. Being able to detect uncertainty in our own knowledge allows us to strategically fill in gaps in understanding by seeking disambiguating information or communicating uncertainty to knowledgeable social partners, or opting out of responding when uncertainty is high. Uncertainty monitoring may be particularly critical early in life, when children are tasked with learning their native language, the social and cultural structure of their world, and the causal properties of their environment. However, there has been mixed evidence about the extent to which young children are conscious of their own knowledge and mental states, and their ability to act on this meta-awareness. This raises questions about how active young children's learning is; do preschool-aged children monitor uncertainty and guide their learning behaviors on the basis of this monitoring, or is early learning better characterized as a process of integrating information that is largely generated externally?

Research on children's uncertainty monitoring within the metacognitive framework has shown that children fail at reporting on their ongoing thoughts (Flavell ref) and are inconsistently able to explicitly report on their confidence in their knowledge. For example, 3-year-olds report being equally sure about correct and incorrect responses in a mem-

ory task (Hembacher ref). There is also protracted development throughout childhood of other metacognitive abilities, such as estimating future performance (JOL refs; Destan ref; Lipowski ref) or selectively studying difficult or less-well-learned materials (). However, most of these studies have relied upon explicit reports of uncertainty or learning progress. It is possible that children are able to act upon their environment and their own knowledge on the basis of uncertainty monitoring prior to their ability to bring it fully into consciousness or organize an explicit response about it.

In support of this possibility, several studies have shown that infants and toddlers engage in spontaneous information-seeking behaviors selectively in response to epistemic uncertainty. For example, Call and Carpenter (2002) had 2-year-olds choose between several tubes to find a hidden sticker. They found that the toddlers were more likely to peek inside a tube before choosing when they had not seen the baiting of the tubes compared to when they had, suggesting they were aware of their ignorance and managed to delay their response until they had a good answer. In another study, Goupil (year) found that 20-month-olds were more likely to seek help by looking at their parents when they were unable to respond accurately in a memory task. These implicit measures may better capture uncertainty signals by bypassing the need to orchestrate an explicit response, even a non-verbal one. Thus, examining spontaneous information-gathering behaviors may resolve remaining questions about early uncertainty monitoring; for example, are young children sensitive to graded uncertainty, or do they merely monitor the presence or absence of knowledge?

Here, we focus on the role of uncertainty in guiding social referencing — one form of information gathering — during word learning. One of the cues children have available to them when learning object-label mappings is the direction of a speaker's gaze, as people tend to look at objects they are referring to (ref). By the second year of life infants follow a speaker's gaze and map labels to objects on the basis of gaze direction (ref). There is also evidence that infant's propensity for gaze-following predicts later language development (ref), highlighting the importance of this behavior for learning.

Although gaze-following is clearly critical for word-object mapping, at any given moment, there are many stimuli in the visual field that compete for attention. In some cases, a listener may be required to disengage from a salient or informative stimulus in order to reference a speaker's gaze direction to resolve referential ambiguity. Thus, it may be optimal to follow a speaker's gaze only when disambiguating information is needed, which would require the listener to monitor

her own uncertainty. Investigating the selectivity of gaze-following and social referencing in general may be a particularly useful case study of uncertainty monitoring processes for several reasons. First of all, social referencing is a ubiquitous spontaneous behavior across the lifespan, and thus does not have to be prompted and is ecologically valid. Second of all, looking is a continuous measure, allowing us to determine whether it, and therefore the monitoring system, is sensitive to graded uncertainty.

Do infants and children follow a speaker's gaze regardless of the need for disambiguation, or do they selectively follow gaze when the referent of a word is unknown? Vaish, Demir and Baldwin (2011) have previously investigated this question with 13- and 18-month-olds infants. Infants sat across from an experimenter who produced a label (e.g., "a modi!") in the presence of one or two objects. They found that infants looked up to the experimenter more often when there were two objects present, and the referent was thus ambiguous. They interpret this to mean that infants recognize when they need disambiguating information and reference gaze accordingly. In addition, when only one object was present, infants successfully mapped the new label to the object, as confirmed in a comprehension test in which children had to place a named object in a bucket. In sum, infants appear to selectively reference a speaker when uncertainty about the referent of a label is high.

The present research adapts Vaish, Demir and Baldwin to examine whether preschool-aged children's social referencing is sensitive to uncertainty based on referential ambiguity, and whether it reflects graded uncertainty about a word-object mapping. In addition to measuring children's looking during the labeling event itself, we measured looking across an event in which children heard a label and then were asked to select the corresponding object and put it in a bucket. We predicted that children might selectively reference their interlocutor on the basis of uncertainty while planning and executing their decision in addition to during the labeling itself, as social information might be expected to be available throughout the mapping and decision processes.

Experiment 1

In Experiment 1, we examined whether children would reference a speaker's face more often when there was referential ambiguity associated with a label she produced. We were interested not only in children's referencing of gaze direction during the labeling itself, but referencing of the speaker after the labeling event while children made a decision about which object was the intended referent. We predicted that children would seek confirmation of the accuracy of their choice from the speaker when they were uncertain, perhaps expecting an emotional reaction or explicit approval, in addition to seeking gaze direction information during labeling. We will refer to looks during and after labeling as social referencing, as the looks could reflect different types of information gathering.

We had children sit across from an experimenter who la-

beled an object on the table between them (Figure 1). Across trials, there were either one or two objects, which were either familiar or unfamiliar to the child. We expected children to be uncertain about the label's referent only on trials with two unfamiliar objects. After labeling an object, the experimenter asked the child to place the named object in a bucket. We measured the number of times the child looked up to the experimenter's face during four discreet phases of the trial: the labeling event (*label* phase), the sliding of the object(s) into the child's reach (*slide* phase), the time before the child touched an object once they were within reach (*planning* phase), and the time between touching an object and dropping it in the bucket (*response* phase).

We predicted that children would look to the experimenter more often on two-object novel trials compared to the other three trial types during the labeling, planning, and response phases, which would indicate that they recognized the need for disambiguating information. We did not predict that children would look more for these trial types during the slide phase, when they would likely be looking at the objects themselves.

Methods

Participants We recruited a planned sample of 80 children ages 2-5 years from the Children's Discovery Museum in San Jose, California. The sample included 20 2-year-olds (mean age 31.71 months), 20 3-year-olds (mean age 42.65 months), 20 4-year-olds (mean age 55.85 months), and 20 5-year-olds (mean age 65.11 months). An additional 20 children participated but were removed from analyses because they heard English less than 75% of the time at home ($n = 10$), because they were unable to complete at least half of the trials in the task ($n = 4$), because of parental interference ($n = 1$), or due to experimenter or technical errors ($n = 5$).

Stimuli and Design In this task, children were presented with one or two objects, heard a label, and were asked to put the labeled object in a bucket. Half of the objects were selected to be familiar to children (e.g., a cow) and half were selected to be novel (e.g., a nozzle). There were four trial types: one-familiar, one-novel, two-familiar, and two-novel. There were three trials of each type, for a total of twelve trials, and trial types were presented sequentially in an order that was counterbalanced across participants. The assignment of individual objects to trial types was counterbalanced. On familiar trials, the familiar label for the target object was used (e.g., "cow"). On novel trials, a novel label was used (e.g., "dawnoo").

The critical manipulation was of referential ambiguity; on one-familiar and two-familiar trials, there was no referential ambiguity, as children were expected to be certain about the objects and their labels. Similarly, on one-novel trials, children were expected to be certain about the label referent as there was only one option. However, on two-novel trials, the referent was ambiguous, as the novel label could apply to either novel object.

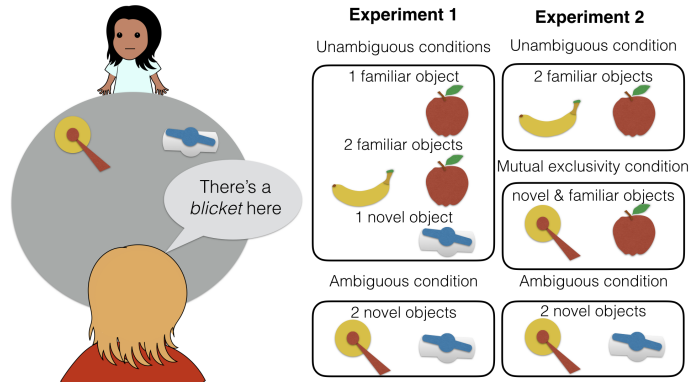


Figure 1: Study design for Experiments 1 and 2.

Procedure Throughout the study, the child sat at one end of a large circular table, and the experimenter stood at the opposite end. Each trial of the task proceeded as follows: the experimenter placed one or two objects on the left and/or right sides of the table, out of reach of the child so that the child could not interact with the toys during the labeling event. For one-object trials, the location of the object (left or right) alternated between trials. After placing the objects, the experimenter said “Hey look, there’s a (target) here.” The experimenter gazed at the center of the table rather than the object she was labeling because we wanted to preserve the referential ambiguity throughout the trial. The experimenter waited approximately two seconds (based on a visual metronome placed within view) before saying, “Can you put the (target) in the bucket?” She then pushed the object(s) forward within reach of the child, and placed a plastic bucket in the center of the table, also within reach of the child. Prior to the twelve experimental trials, there were two training trials: a one-familiar trial and a two-familiar trial, to acquaint the child with the procedure. A camera placed to the side of the experimenter captured the participant’s face, so that looking behavior could be coded from video.

Coding procedure Videos were coded using DataVyu software (<http://datavyu.org>). First, each trial was divided into four temporal phases: a *label* phase, which began at the utterance of the target label and ended when the experimenter began to slide the objects, a *slide* phase, which encompassed the sliding of the objects into the child’s reach, a *planning* phase, which began at the end of the slide and ended when the child touched an object, and a *response* phase, which began when the child touched an object and ended when the child released the object into the bucket. After onsets and offsets of these phases had been coded, the coder recorded the number of looks the child made to the experimenter’s face during each phase. We opted to code the number of looks rather than the duration of looks because we felt that looks from the stimuli to the experimenter’s face and vice versa might allow children to integrate social and nonsocial information to solve the problem of reference (refs), and might also be less sensitive

to individual differences in the duration of individual looks. A second coder coded the number of looks for a quarter of the trials for each participant to establish reliability. *include reliability data here*

Results and Discussion

Descriptive results of Experiment 1 are presented in Figure 2. To test our prediction that referential ambiguity (i.e., having two novel objects) would produce more social referencing, we fit mixed-effects linear regression models separately for each phase with the following structure: `number of looks ~ number of objects * familiarity * age in months + (number of objects + familiarity | SID)`. A single model with phase as a factor did not converge.

We did not find any main or interactive effects of number of objects, familiarity, or age on number of looks during the label phase or the slide phase. However, we found an interactive effect of number of objects and familiarity during the planning ($\beta = 0.21, p < .001$) and response phases ($\beta = 0.6, p < .001$), such that 2-novel trials were associated with more looking. There was no interaction with age in either phase¹. In summary, children ages 2-5 looked to the experimenter more often when planning and executing a response under uncertainty. These results suggest that children were aware that they did not have sufficient knowledge to answer independently, and they attempted to resolve their uncertainty using social referencing.

We did not find the expected effect of referential ambiguity in the label phase. There are a number of reasons we might have observed this null effect. One possibility is that children failed to predict that they would need more information until later in the trial, when they were actually faced with planning and executing a decision. Another possibility is that children’s looking was at ceiling during the labeling phase, perhaps because children look at someone who is speaking regardless of the need for referential disambiguation. A third possibility is that this is an artifact of our design, in which the experimenter gazed at the center of the table rather than the

¹Code and data available with full regression model results at <https://github.com>

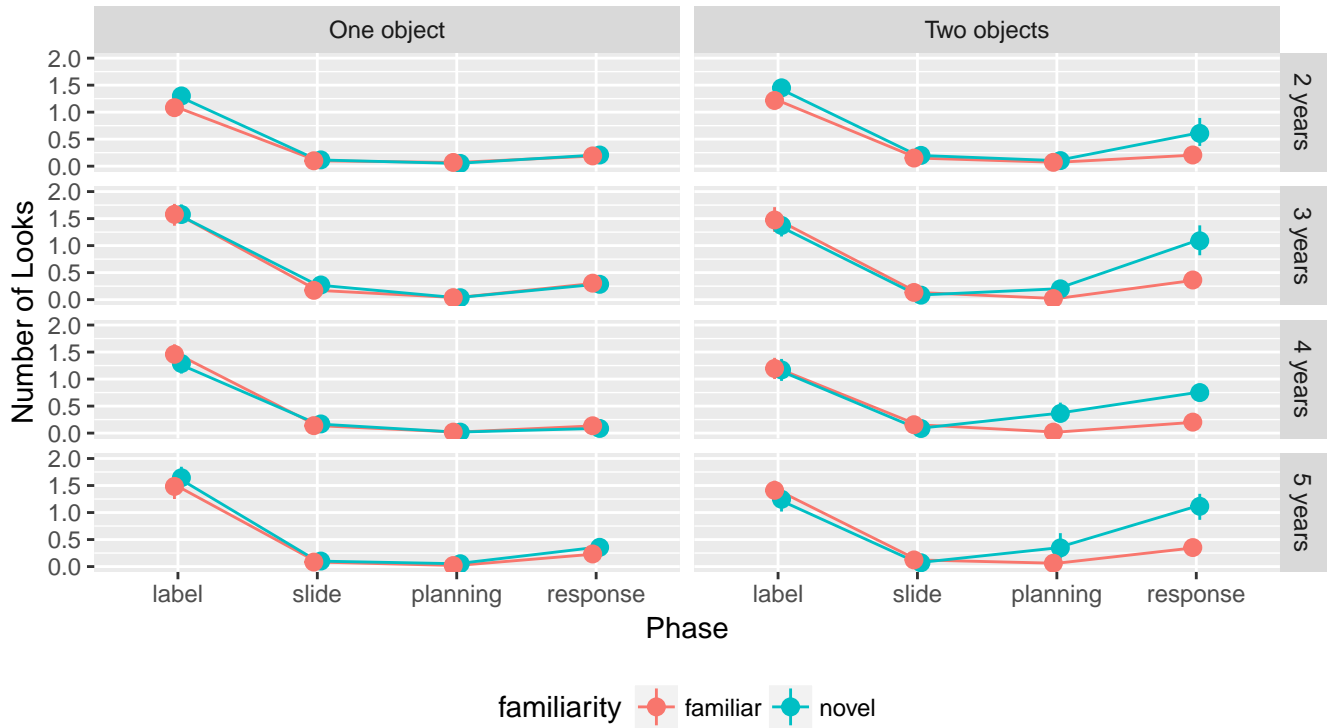


Figure 2: Results of Experiment 1. Number of looks to the experimenter across phases and conditions. Error bars are 95 percent confidence intervals.

referent of her label. Children may have realized that the experimenter’s gaze direction during labeling was not a source of disambiguating information. Experiment 2 seeks to examine this possibility.

Experiment 2

Experiment 2 was designed to replicate Experiment 1 and address the possibility that the experimenter’s gaze pattern during labeling had an effect on children’s social referencing. To test this possibility, we manipulated the experimenter’s gaze behavior between participants. For half of the sample, the experimenter gazed at the center of the table during labeling (i.e., reproducing Experiment 1), and for the remaining half, the experimenter looked at the object she was referring to during labeling. In Experiment 1, we did not observe an effect of age on looking. Thus, we restricted the current sample to 3- and 4-year-olds. Finally, since we did not observe any difference between the one-familiar and one-novel trials, we eliminated single-object trials, leaving the 2-familiar and 2-novel trials. In addition to these trials, we added 1-novel-1-familiar trials to examine children’s looking patterns when 2 objects are present and an unfamiliar label are produced, but the label referent can be deduced through the principle of mutual exclusivity.

Methods

Participants We recruited a planned sample of 80 children ages 3-4 years from the Children’s Discovery Museum in

San Jose, California. The sample included 40 3-year-olds (mean age 42.89 months) and 40 4-year-olds (mean age 53.47 months). An additional 20 children participated but were removed from analyses because they heard English less than 75% of the time at home ($n = 9$), because they were unable to complete at least half of the trials in the task ($n = 7$), or due to experimenter or technical errors ($n = 4$).

Stimuli and Design The stimuli and design were similar to Experiment 1, except that we eliminated 1-object trials. Instead, we included three trial types: 2-familiar (“familiar”), 2-novel (“novel”), and 1-novel-1-familiar (“mutual exclusivity”). There were four of each trial type, totaling twelve trials. In addition, we manipulated the experimenter’s gaze behavior between participants. For half of the participants, she looked at the center of the table on every trial; for the remaining half, she looked at the object she referred to on every trial.

Procedure The procedure was identical to Experiment 1, except that there were three practice trials rather than two, so that children could experience every trial type.

Results and Discussion

The descriptive results of Experiment 2 are presented in Figure 3. To quantify the main and interactive effects of familiarity, number of objects, gaze behavior, and age on social referencing, we fit a mixed-effects linear regression model with the following structure: `number of looks ~ familiarity * age`




Figure 3: Results of Experiment 2. Number of looks to the experimenter across phases and conditions. Error bars are 95 percent confidence intervals.

in months * gaze * phase + (familiarity | SID).

As in Experiment 1, there was an interaction of phase with familiarity such that the response phase of novel trials was associated with significantly more looks ($\beta = 0.51, p < .001$). There was a three-way interaction of familiarity, gaze, and phase, such that the response phase of mutual exclusivity trials in the no-gaze condition was associated with more looks ($\beta = 0.39, p < .01$). Finally, we observed a four-way interaction such that the response phase of novel trials in the gaze condition was associated with more looking with increasing age ($\beta = 0.06, p < .01$).

Overall, these results replicate the finding from Experiment 1 that children engage in more social referencing during the response phase (i.e., when executing a decision about which item is being referred to) when there is referential ambiguity because two novel objects are present. However, we did not replicate the finding that children engaged in more social referencing during the planning phase on novel trials. At a descriptive level there was a trend for novel trials to be associated with more looking in the no-gaze condition but not in the gaze condition, perhaps because children were less uncertain after having received helpful referential gaze during labeling.

In addition, we did not find selective social referencing during the label phase, even when referential gaze was available. This rules out the possibility that children were less selective during this phase because they realized that gaze direction information was not available. Instead, they might be at ceiling for looking (indeed, they do more looking during this phase than the other three phases, regardless of condition), or they may not recognize the additional need for information in novel trials until they are faced with making a decision.

Interestingly, the three-way interaction of familiarity, gaze, and phase showed that mutual exclusivity trials were asso-

ciated with greater looking during the response phase when the experimenter gazed at the object she labeled compared to when she did not. This is intriguing given that children should be able to solve mutual exclusivity trials without gaze information. Instead, it seems that they remain somewhat uncertain while executing a decision, but this uncertainty is resolved if the experimenter gazes at the correct object during labeling.

Finally, the four-way interaction with age suggests that children selectively reference social information in response to uncertainty to a greater degree as they get older. This could be due to multiple factors; on the one hand, they may more accurately monitor the need for disambiguating information with development. On the other hand, they may be more likely to recognize that social information can be a source of disambiguation with age. Regardless, the interaction with age should be interpreted with caution as we did not find any interaction with age in Experiment 1.

General Discussion

Preschoolers quickly learn new concepts, rules, and language, often with minimal exposure (refs). They also actively explore and ask questions in ways that seem targeted to maximize learning (Choinard, Schulz & Bonawitz). However, we still have an incomplete understanding of young children's ability to monitor their own mental states, in particular, their epistemic uncertainty. Do children monitor their own uncertainty to guide information seeking behaviors, or are external features of the environment sufficient to guide these behaviors (e.g., children might ask for help with a complex looking toy in response to perceptual features rather than epistemic states). Here, we examined a spontaneous information-seeking behavior, social referencing, and its selectivity with regard to epistemic uncertainty among preschool-aged chil-

dren. Specifically, we asked whether children would reference a social partner more often when there was referential ambiguity in a label she produced, and whether the amount of social referencing would reflect graded uncertainty based on the availability of cues to reference.

We found strong evidence of selective social referencing under referential ambiguity during the response phase of the trial, after children had touched one of the objects but had not yet committed to it by placing it in the bucket. In Experiment 1, we additionally found evidence for selective social referencing earlier in the trial, after the objects had been placed within reach but before the child touched one of them, though this was not replicated in Experiment 2. Although we cannot determine precisely what information children hoped to obtain through these looks, we speculate that they may have expected confirmation of the accuracy of their choice, either implicitly through the adult's facial expressions, or through explicit feedback.

Importantly, we also found evidence for selectivity in social referencing based on graded uncertainty. In Experiment 2, we manipulated the amount of evidence available to children by including three trial types: familiar trials in which children were familiar with the label-object pairing and uncertainty should be low, novel trials in which children were unfamiliar with the object-label pairing and uncertainty should be high, and trials in which children could use mutual exclusivity to discover the label-object pairing, but were not previously familiar with the label or object. We also manipulated whether or not gaze direction during labeling was useful; half of participants had an experimenter who gazed at the object she labeled, and half had one who gazed at neither object and was thus uninformative.

We found that children treated mutual exclusivity trials more like familiar trials in terms of the amount of looking, but only when they had received helpful gaze, suggesting that the combination of mutual exclusivity cues and gaze direction cues were required for children to feel confident about the object-label mapping. When mutual exclusivity trials were not paired with helpful gaze, they were treated more like novel trials. Importantly, useful gaze during labeling did not lessen the amount of social referencing during the response phase for novel trials, suggesting that gaze information alone was not sufficient to reduce uncertainty. This means that children demonstrate uncertainty in the presence of multiple objects and an unfamiliar label even when they should be able to answer independently. It is possible that they remain uncertain about a new label even after they have acquired it, if they have only heard it once and not received confirmation of its accuracy, for example, through gaze monitoring.

On the other hand, we found no evidence for selectivity as the object was being labeled, or as the objects were being slid into reach. One interpretation of this pattern of results is that preschool-aged children do not recognize the need for disambiguating information when a referent is ambiguous until they are in the position of needing to make a decision, at the end of

the trial. However, another possibility is that children spontaneously look at someone who is speaking regardless of the need for disambiguating information, and additional looking on top of this baseline was not needed or possible. A related possibility is that the labeling phase of our task was too fast for children to engage in extra looking on top of baseline in response to uncertainty. Note that Vaish, Demir and Baldwin found that infants selectively referenced a speaker during the labeling itself. However, their procedure was different from the present study in that children were holding one of the toys as the speaker produced a label. Thus, referencing the speaker required disengaging from an interesting toy that they were already attending to. This may have prevented ceiling levels of looking and forced infants to be selective. A follow-up to the present work that includes a greater trade-off in attentional possibilities would help to distinguish among these possibilities.

Overall, these results confirm that preschool-aged children monitor graded uncertainty in their knowledge and act on that uncertainty through information-gathering behaviors. These findings contribute to an emerging view of children's early learning as active and driven by probabilistic evidence.

Acknowledgements

We thank Veronica Cristiano for assisting with data collection.

References