Children's social referencing reflects sensitivity to graded uncertainty

Emily Hembacher
ehembach@stanford.edu
Department of Psychology
Stanford University

Benjamin deMayo
bedemayo@stanford.edu
Department of Psychology
Stanford University

Michael C. Frank
mcfrank@stanford.edu
Department of Psychology
Stanford University

Abstract

abstract abstract.

Keywords: social referencing; help seeking; word learning; uncertainty.

There has been a great deal of debate over the extent to which young children are aware of their mental states, including epistemic uncertainty and ignorance. Although it was traditionally assumed that preschool-aged children were either unaware of their possible ignorance or underestimated it (Flavell, Schneider?), newer evidence shows that preschoolers can explicitly report on their uncertainty in their knowledge, and infants and toddlers react in ways that demonstrate sensitivity to uncertainty. For example, preschool-aged children can use a confidence scale to accurately report that they are less sure about inaccurate memory (Hembacher & Ghetti, 2014) or perceptual discrimination decisions (Lyons & Ghetti, Coughlin, Hembacher, Lyons & Ghetti).

Other research has investigated the spontaneous behaviors that infants and toddlers engage in in response to epistemic uncertainty. For example, Call and Carpenter (2002) had 2-year-olds choose between several tubes to find a hidden sticker. They found that the toddlers were more likely to peek inside a tube before choosing when they had not seen the baiting of the tubes compared to when they had. In another study, Goupil (year) found that 20-month-olds were more likely to seek help by looking at their parents when they were unable to respond accurately in a memory task.

This evidence suggests that humans are able to compute the likely accuracy of their decisions or beliefs and act accordingly from a young age. However, even older children have been found to fail at metacognitive tasks that appear to rely on uncertainty monitoring. For example, Hembacher & Ghetti (2014) found that 3-year-olds reported being equally confident for correct and incorrect memory decisions, and there is protracted development of other metacognitive judgments of performance, such as judgments of how well material has been learned, through adolescence (refs).

Given these somewhat contradictory findings, it might be that the ability to detect or estimate uncertainty emerges early in development, but the ability to organize metacognitive thought processes about these computations and translate them into strategic behavior develops more slowly. Similar arguments have been made with regard to the distinction between metacognitive monitoring and control; in the metacognitive theoretical framework proposed by Nelson and Narens (1990), metacognitive monitoring involves awareness of mental states, and metacognitive control involves acting upon these states, for example, by seeking new information to resolve gaps in knowledge. Many studies have found that children's ability to report on the likely accuracy of their knowledge or performance developmentally precedes their ability to act on this awareness, for example, by spending more time learning difficult compared to easy materials (refs) or refusing to answer when uncertain (refs). Young children may react to uncertainty by hesitating or seeking help when faced with a difficult task (refs), while older children become increasingly able to coordinate more sophisticated learning behaviors in response to uncertainty (refs).

Despite this progress in characterizing children's uncertainty monitoring and related behaviors, there is still little known about the types and sources of uncertainty that children are able to monitor, and what spontaneous behaviors they may engage in in response to this monitoring. These issues tie into larger questions about children's learning. A great deal of research has focused on characterizing the ways that young children make sense of input from their environment, for example, when acquiring language or concepts, but there is increasing interest in the role that children play in orchestrating the flow of information itself (refs). Uncertainty monitoring may be a critical factor in children's active information gathering behaviors.

One possible instantiation of this interplay may be in children's gaze following during word learning. Gaze following allows individuals to infer the likely referent of novel labels that occur in a speech stream, as speakers typically gaze at their intended referent (refs). However, at any given moment, there are multiple stimuli in the visual field that compete for attention. In some cases, a listener may be required to disengage from a salient or informative stimulus in order to reference a speaker's gaze direction to resolve referential ambiguity. Thus, an optimal system might involve referencing a speaker's gaze only when referential ambiguity is high and gaze direction information is necessary for disambiguation. Vaish, Demir and Baldwin (2010) found evidence to support this possibility among 12-18-month-old infants. They had infants sit across from an experimenter with either one or two objects present. Infants were more likely to look up at the speaker when she produced a novel label when there were two objects, which produced referential ambiguity, suggesting that they were sensitive to the need for disambiguating information.

The present research uses a similar approach to investigate preschool-aged children's sensitivity to uncertainty based on referential ambiguity. We aimed to replicate Vaish, Demir and Baldwin in a new age group, and investigate the time

course of social referencing during an event in which children hear a label under ambiguous or unambiguous conditions and are asked to select the corresponding object. This procedure establishes social referencing as an implicit and spontaneous measure of uncertainty monitoring, which could potentially be used across a broad age range to investigate the development of uncertainty monitoring and corresponding behavior.

Experiment 1

In Experiment 1, we examined whether children would reference a speaker's face more often when there was referential ambiguity associated with a label she produced. We were interested not only in children's referencing of gaze direction during the labeling itself, but referencing of the speaker after the labeling event while children made a decision about which object was the intended referent. We predicted that children would seek confirmation of the accuracy of their choice from the speaker when they were uncertain, perhaps expecting an emotional reaction or explicit approval, in addition to seeking gaze direction information during labeling. We will refer to all of this looking behavior as social referencing, as the looks could reflect different types of information gathering.

To test this, we had children sit across from an experimenter who labeled an object on the table between them. Across trials, there were either one or two objects, which were either familiar or unfamiliar to the child. We expected children to be uncertain about the label's referent only on trials with two unfamiliar objects. After labeling an object, the experimenter asked the child to place the named object in a bucket. We measured the number of times the child looked up to the experimenter's face during four discreet phases of the trial: the labeling event ("label" phase), the sliding of the object(s) into the child's reach ("slide" phase), the time before the child touched an object once they were within reach ("planning" phase), and the time between touching an object and dropping it in the bucket ("response" phase).

We predicted that children would look to the experimenter more often on two-object novel trials compared to the other three trial types during the labeling, planning, and response phases, which would indicate that they recognized the need for disambiguating information. We did not predict that children would look more for these trial types during the slide phase, when they would likely be looking at the objects themselves.

Methods

Participants   We recruited a planned sample of 80 children ages 2-5 years from the Children's Discovery Museum in San Jose, California. The sample included 20 2-year-olds (mean age 31.71 months), 20 3-year-olds (mean age 42.65 months), 20 4-year-olds (mean age 55.85 months), and 20 5-year-olds (mean age 65.11 months). An additional 20 children participated but were removed from analyses because they heard English less than 75% of the time at home ($n = 10$), because they were unable to complete at least half of the trials in the

task ($n = 4$), because of parental interference ($n = 1$), or due to experimenter or technical errors ($n = 5$).

Stimuli and Design   In this task, children were presented with one or two objects, heard a label, and were asked to put the labeled object in a bucket. Half of the objects were selected to be familiar to children (e.g., a cow) and half were selected to be novel (e.g., a nozzle). There were four trial types: one-familiar, one-novel, two-familiar, and two-novel. There were three of each trial type, for a total of twelve trials, and trial types were presented sequentially in an order that was counterbalanced across participants. The assignment of individual objects to trial types was counterbalanced. On familiar trials, the familiar label for the target object was used (e.g., "cow"). On novel trials, a novel label was used (e.g., "dawnoo").

The critical manipulation was of referential ambiguity; on one-familiar and two-familiar trials, there was no referential ambiguity, as children were expected to be certain about the objects and their labels. Similarly, on one-novel trials, children were expected to be certain about the label referent as there was only one option. However, on two-novel trials, the referent was ambiguous, as the novel label could apply to either novel object (Figure 1).

Procedure   Throughout the study, the child sat at one end of a large circular table, and the experimenter stood at the opposite end. Each trial of the task proceeded as follows: the experimenter placed one or two objects on the left and/or right sides of the table, out of reach of the child so that the child could not interact with the toys during the labeling event. For one-object trials, the location of the object (left or right) alternated between trials. After placing the objects, the experimenter gazed at the center of the table and said "Hey look, there's a (target) here." The experimenter waited approximately two seconds (based on a visual metronome placed within view) before saying, "Can you put the (target) in the bucket?" She then pushed the object(s) forward within reach of the child, and placed a plastic bucket in the center of the table, also within reach of the child. Prior to the twelve experimental trials, there were two training trials: a one-familiar trial and a two-familiar trial, to acquaint the child with the procedure. A camera placed to the side of the experimenter captured the participant's face, so that looking behavior could be coded from video.

Coding procedure   Videos were coded using DataVyu software (http://datavyu.org). First, each trial was divided into four temporal phases: a "label" phase, which began at the utterance of the target label and ended when the experimenter began to slide the objects, a "slide" phase, which encompassed the sliding of the objects into the child's reach, a "planning" phase, which began at the end of the slide and ended when the child touched an object, and a "response" phase, which began when the child touched an object and ended when the child released the object into the bucket. After onsets and offsets of these phases had been coded, the
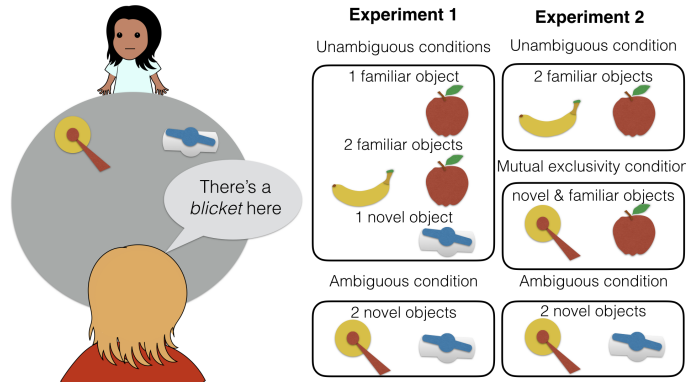
Figure 1: Study design for Experiments 1 and 2.

coder recorded the number of looks the child made to the experimenter's face during each phase. A second coder coded the number of looks for of the trials for each participant to establish reliability. *include reliability data here*

Results and Discussion

The descriptive results of Experiment 1 are presented in Figure 2. To quantify the main and interactive effects of referential ambiguity and age on looking, we fit mixed-effects linear regression models separately for each phase with the following structure: `number of looks ~ number of objects * familiarity * age in months + (number of objects + familiarity | SID)`. A single model with phase as a factor did not converge.

We did not find any main or interactive effects of number of objects, familiarity, or age on number of looks during the label phase or the slide phase. However, we found an interactive effect of number of objects and familiarity during the planning ($\beta = 0.21$, $p < .001$) and response phases ($\beta = 0.6$, $p < .001$), such that 2-novel trials were associated with more looking. There was no interaction with age in either phase.

In summary, we found that children ages 2-5 looked to the experimenter more often when planning and executing a response under uncertainty. This suggests that they were aware that they did not have sufficient knowledge to answer independently, and they attempted to resolve their uncertainty using social referencing. However, we did not find the expected effect of referential ambiguity in the label phase. There are a number of reasons we might have observed this null effect. One possibility is that children failed to predict that they would need more information until later in the trial, when they were actually faced with planning and executing a decision. Another possibility is that children's looking was at ceiling during the labeling phase, perhaps because children look at someone who is speaking regardless of the need for referential disambiguation. A third possibility is that this is an artifact of our design, in which the experimenter gazed at the center of the table rather than the referent of her label. Children may have realized that the experimenter's gaze direction during labeling was not a source of disambiguating

information. Experiment 2 seeks to examine this possibility.

Experiment 2

Experiment 2 was designed to replicate Experiment 1 and address the possibility that the experimenter's gaze pattern during labeling had an effect on children's social referencing. To test this possibility, we manipulated the experimenter's gaze behavior between participants. For half of the sample, the experimenter gazed at the center of the table during labeling (i.e., reproducing Experiment 1), and for the remaining half, the experimenter looked at the object she was referring to during labeling. In Experiment 1, we did not observe an effect of age on looking. Thus, we restricted the current sample to 3- and 4-year-olds. Finally, since we did not observe any difference between the one-familiar and one-novel trials, we eliminated single-object trials, leaving the 2-familiar and 2-novel trials. In addition to these trials, we added 1-novel-1-familiar trials to examine children's looking patterns when 2 objects are present and an unfamiliar label are produced, but the label referent can be deduced through the principle of mutual exclusivity.

Methods

Participants We recruited a planned sample of 80 children ages 3-4 years from the Children's Discovery Museum in San Jose, California. The sample included 40 3-year-olds (mean age 42.89 months) and 40 4-year-olds (mean age 53.47 months). An additional 20 children participated but were removed from analyses because they heard English less than 75% of the time at home ($n = 9$), because they were unable to complete at least half of the trials in the task ($n = 7$), or due to experimenter or technical errors ($n = 4$).

Stimuli and Design The stimuli and design were similar to Experiment 1, except that we eliminated 1-object trials. Instead, we included three trial types: 2-familiar ("familiar"), 2-novel ("novel"), and 1-novel-1-familiar ("mutual exclusivity"). There were four of each trial type, totaling twelve trials. In addition, we manipulated the experimenter's gaze behavior between participants. For half of the participants, she looked
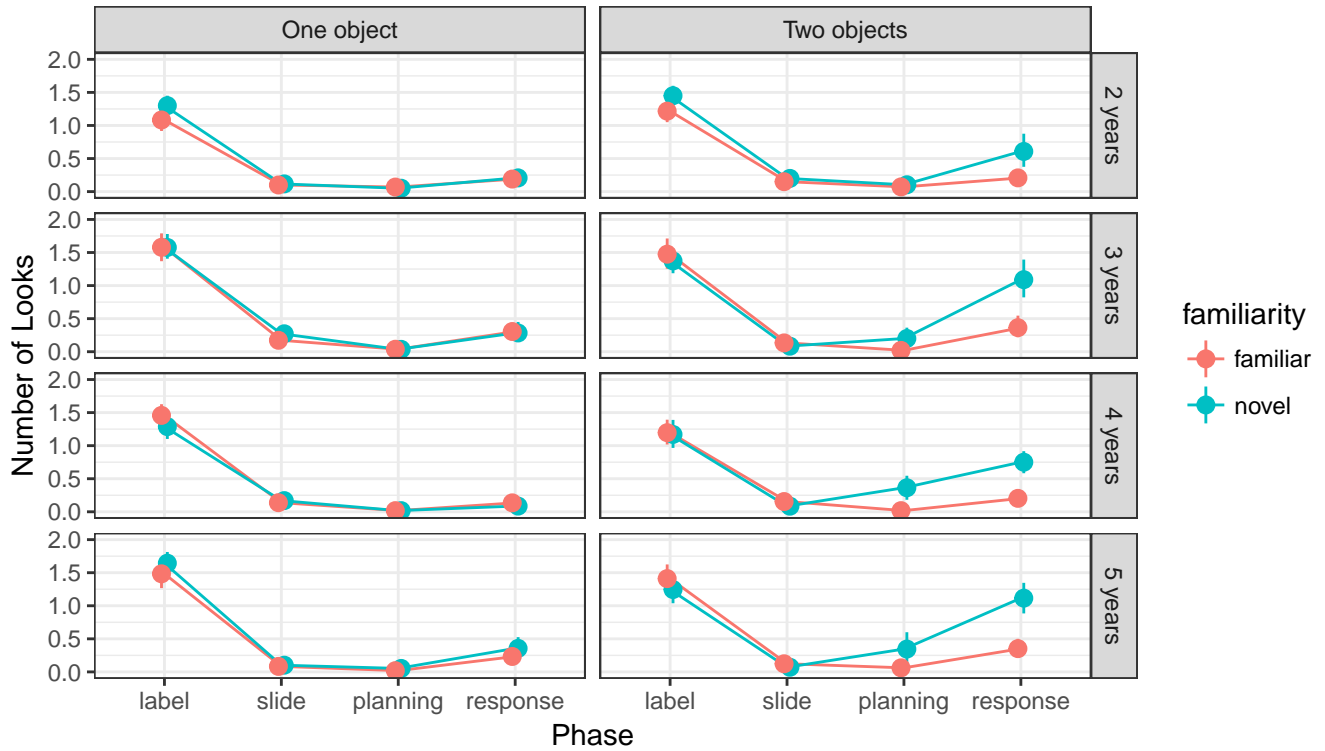
Figure 2: Results of Experiment 1. Number of looks to the experimenter across phases and conditions.

at the center of the table on every trial; for the remaining half, she looked at the object she referred to on every trial.

Procedure The procedure was identical to Experiment 1, except that there were three practice trials rather than two, so that children could experience every trial type.

Results and Discussion

The descriptive results of Experiment 2 are presented in Figure 3. To quantify the main and interactive effects of familiarity, number of objects, gaze behavior, and age on social referencing, we fit a mixed-effects linear regression model with the following structure: `number of looks ~ familiarity * age in months * gaze * phase + (familiarity | SID)`.

As in Experiment 1, there was an interaction of phase with familiarity such that the response phase of novel trials was associated with significantly more looks ($\beta = 0.51$, $p < .001$). There was a three-way interaction of familiarity, gaze, and phase, such that the response phase of mutual exclusivity trials in the no-gaze condition was associated with more looks ($\beta = 0.39$, $p < .01$). Finally, we observed a four-way interaction such that the response phase of novel trials in the gaze condition was associated with more looking with increasing age ($\beta = 0.06$, $p < .01$).

Overall, these results replicate the finding from Experiment 1 that children engage in more social referencing during the response phase (i.e., when executing a decision about which item is being referred to) when there is referential ambigu-

ity because two novel objects are present. However, we did not replicate the finding that children engaged in more social referencing during the planning phase on novel trials. At a descriptive level there was a trend for novel trials to be associated with more looking in the no-gaze condition but not in the gaze condition, perhaps because children were less uncertain after having received helpful referential gaze during labeling.

In addition, we did not find selective social referencing during the label phase, even when referential gaze was available. This rules out the possibility that children were less selective during this phase because they realized that gaze direction information was not available. Instead, they might be at ceiling for looking (indeed, they do more looking during this phase than the other three phases, regardless of condition), or they may not recognize the additional need for information in novel trials until they are faced with making a decision.

Interestingly, the three-way interaction of familiarity, gaze, and phase showed that mutual exclusivity trials were associated with greater looking during the response phase when the experimenter gazed at the object she labeled compared to when she did not. This is intriguing given that children should be able to solve mutual exclusivity trials without gaze information. Instead, it seems that they remain somewhat uncertain while executing a decision, but this uncertainty is resolved if the experimenter gazes at the correct object during labeling.

Finally, the four-way interaction with age suggests that children selectively reference social information in response
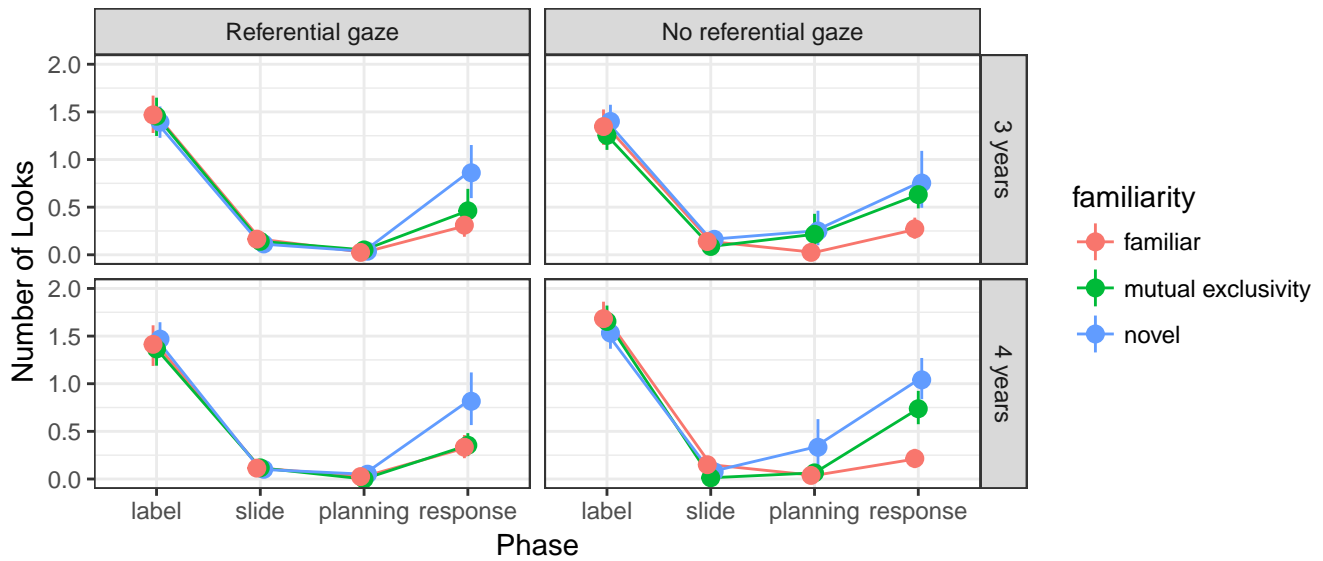
Figure 3: Results of Experiment 2. Number of looks to the experimenter across phases and conditions.

to uncertainty to a greater degree as they get older. This could be due to multiple factors; on the one hand, they may more accurately monitor the need for disambiguating information with development. On the other hand, they may be more likely to recognize that social information can be a a source of disambiguation with age. Regardless, the interaction with age should be interpreted with caution as we did not find any interaction with age in Experiment 1.

General Discussion

The present research examined a spontaneous information-seeking behavior, social referencing, and its selectivity with regard to uncertainty among preschool-aged children. Specifically, we asked whether children would look to an adult's face more often when there was referential ambiguity in a label she produced. Rather than focus exclusively on gaze following, we examined children's social referencing across the timecourse of an event in which children had to decide which object the adult referred to. Although children might have expected different information to be available during different phases of the trial (e.g., gaze direction information while the adult produced a label, and confirmation of accuracy after touching one of the objects), the specific information-seeking goals driving children's social referencing at different times are unknown to us.

Overall, we found strong evidence of selective social referencing under referential ambiguity at the end of the trial, after children had touched one of the objects but had not yet committed to it by placing it in the bucket. In Experiment 1, we found evidence for selective social referencing earlier in the trial, after the objects had been placed within reach but before the child touched one of them, though this was not replicated in Experiment 2. We found no evidence for selectivity at the beginning of the trial, as the object was being labeled, or as the objects were being slid into reach.

One interpretation of this pattern of results is that preschool-aged children do not recognize the need for disambiguating information when a referent is ambiguous until they are in the position of needing to make a decision, at the end of the trial. However, another possibility is that children spontaneously look at someone who is speaking regardless of the need for disambiguating information, and additional looking on top of this baseline was not needed or possible. Note that Vaish, Demir and Baldwin found that infants selectively referenced a speaker during the labeling itself. However, their procedure was different from the present one in that children were holding one of the toys as the speaker produced a label. Thus, referencing the speaker required disengaging from an interesting toy that they were already attending to. This may have prevented ceiling levels of looking and forced infants to be selective. A follow-up to the present work that includes more of a trade-off in attentional options would help to tease apart these possibilities.

We also examined children's social referencing when the referent of a label could be inferred through mutual exclusivity. Children are able to learn new words based on the mutual exclusivity principle before age 2 (Markman, Wasow, & Hansen, 2003). However, when the experimenter did not gaze at the target during labeling, children's referencing of the experimenter on mutual exclusivity trials is similar to novel trials. Thus, children demonstrate uncertainty in the presence of multiple objects and an unfamiliar label even when they should be able to answer independently. It is possible that they remain uncertain about a new label even after they have acquired it, if they have only heard it once and not received confirmation of it's accuracy, for example, through gaze monitoring.

Overall, these results confirm that preschool-aged children are able to monitor their uncertainty in their knowledge and act on that uncertainty through information-gathering behav-

iors. This adds to a growing body of evidence that young children are more metacognitively skilled than previously believed. In addition, this research establishes a new paradigm for measuring children's uncertainty via spontaneous behaviors, which may be a useful complement to research based in the metacognitive framework, which typically involves explicit reports of uncertainty after discrete decisions. The present method of measuring looking behavior could be extended to other cognitive tasks, such as category learning or memory. This could potentially also be used as a proxy for the strength of children's representations such as category boundaries, memories or percepts. In sum, selective social referencing appears to be a promising index of children's uncertainty monitoring.

Acknowledgements

References